

Block cipher modes of operation

Introduction

Since studying block ciphers like AES and DES in the classroom, we neglected the fact that they can only take a specific size of input to encrypt. Upon learning this, we immediately thought “well what should we do if we had a longer message or maybe a shorter one?”, then immediately came to think about the classic first solution to go ahead about it: keep encrypting blocks of data until you reach the last block, and pad it to fill any remainings. What we’ve just described is ECB (Electronic code block) Mode. The simplest (and weakest, as we will see later) block cipher mode of operation.

After doing a little bit more research into the topic, we’ve discovered that there exist many modes of operation; ie. methods that make a plaintext of numerous blocks in size fit into an encryption cipher that only operates on a single block. These modes can be split into two groups: modes that (on their own) only provide confidentiality, others also provide authenticity.

We have also learnt through research that block cipher modes are the weak link in the implementation of block cipher algorithms like AES, not the algorithm itself. AES is considered to be very secure and the best (classical) attempt we have so far only managed to reduce the key size by a couple of bits (in the case of AES-128), with quantum implementations only cutting the key length by half (using AES-256 therefore becomes equivalent to using AES-128). But in this case, it’s not the core block cipher algorithm that is the feasible target, it’s more the way we use to make the plaintext suitable for processing with the block cipher. These block cipher modes of operation remain the most prone to vulnerabilities and exploits of inner workings, and therefore undermine the security of the block cipher algorithm.

Padding

We will be using PKCS#7 to pad plaintexts to fit the block size. It consists of bytes that are added to the plaintext to fill up the last block so that it is the same size as the other blocks. The value of each added byte is equal to the number of bytes added. For example, if one byte is added, then its value is 0x01; if two bytes are added, then their values are 0x02 0x02, and so on. The beauty of this padding scheme is that the receiver (supposing the padding is valid) only needs to look at the last byte and remove that much bytes to get the intended plaintext.

Well known block cipher modes of operation

We attempt to explain the inner workings of some of the most well known modes of operation: ##### ECB ECB (Electronic Codebook) mode is a block-based encryption algorithm that divides the plaintext message into blocks of a fixed size and encrypts each block independently using the same encryption key. The

ciphertext blocks generated by the encryption process are then concatenated to produce the final ciphertext. ECB mode does not introduce any randomness or dependencies between blocks, which makes it vulnerable to certain attacks, including replay attacks and frequency analysis attacks. ECB mode is parallelizable for both encryption and decryption, meaning that multiple blocks can be encrypted or decrypted simultaneously, without needing to wait for the previous block to complete. This property makes it pretty efficient and reduces processing times. ECB mode also allows for random read access – the ability to access any block of ciphertext directly, without needing to decrypt the previous blocks in the message – as each block is encrypted independently.

CBC CBC (Cipher Block Chaining) mode, on the other hand, introduces an additional level of randomness and interdependence between blocks. In CBC mode, the plaintext message is XORed with the previous ciphertext block before being encrypted. This process is repeated for each block in the message, with the first block being XORed with a random initialization vector (IV). CBC is more secure than ECB because identical plaintext blocks will not produce identical ciphertext blocks, as each block depends on the previous block's ciphertext. CBC is parallelizable for encryption but not for decryption, it also doesn't provide random read access: in order to decrypt a block, we need the previous block.

OFB OFB (Output Feedback) mode is a special type of block cipher mode of operation, because it turns the block cipher into a stream cipher. Consequently, it eliminates the need for padding schemes, and produces a ciphertext the length of the original message. OFB generates a keystream of random bits, which is then XORed with the plaintext to produce the ciphertext. The first keystream block is generated independently of the plaintext by encrypting a fixed initialization vector (IV) with the encryption key. This keystream block is also fed back into the encryption process to generate the next keystream block, which is used to encrypt the second plaintext block, and so on. One advantage of OFB mode is that it is a self-synchronizing cipher, meaning that errors in transmission or encryption do not propagate to affect other parts of the message. This is because the keystream is generated independently of the plaintext and is not affected by errors in the transmission of the ciphertext. Another advantage of OFB mode is that it allows for random read access, which means that any block of ciphertext can be decrypted without requiring the previous blocks. This is because the keystream is generated independently of the plaintext, and each block of the keystream is only dependent on the encryption key and the IV. However, one potential weakness of OFB mode is that if the keystream is ever reused with the same IV and encryption key, then an attacker can recover the plaintext by XORing the ciphertext with the keystream. Therefore, it is essential to use a unique IV for each encryption, and to ensure that the IV is not reused with the same encryption key.

CTR CTR (Counter) mode works like OFB: it turns a block cipher into a stream cipher. But instead of using an IV to introduce randomness, it generates the next keystream block by encrypting successive values of a concatenation between a nonce and a counter. A nonce can be thought of just like an IV: a sequence of random bytes with a size half of the block size, the other half is filled with increments of the counter.

One important feature of both the CTR and the OFB modes is that the encryption function is in itself the decryption function. No need to implement two separate algorithms, since they're basically just XORing a keystream with the plaintext to produce ciphertext or vice-versa.

Attack

For demonstrating attacks, and because of the difficulty of the implementation of these attacks, we've only chosen to concentrate on four different attacks on two block cipher modes of operation: ECB and CBC.

Encryption Oracle

An encryption oracle is just a function which can tell us something about a given input we feed to it. It usually returns information in the form of a boolean value (true/false), this in return provides us with information to let us carry on our attack.

Feasibility of Chosen Plaintext Attacks (CPA)

Chosen plaintext attacks is a cryptanalysis model which presumes that the attacker can obtain the ciphertext for arbitrary plaintext; he basically has control over (part of) the plaintext that's being encrypted. We start by revealing why this assumption isn't that far fetched, since we will be demonstrating some attacks based on it later. A pretty demonstrating example of this is the web pages we use today that run on plenty of javascript. Having the target redirected to a website owned by the attacker for example could lead to the website injecting some malicious javascript code that, while having limited capabilities, might still have the ability to perform a CPA.

Detecting the Block Size

We have to start gathering as much informations as we can on the plaintext, the first and easiest thing to start with would be the block size. And to know that, we will use a simple scheme which consists of providing longer and longer input to the encryption function until another block is added. Then we just calculate the difference in length to get the block size used in the encryption oracle.

ECB

Detecting ECB Mode An example function which can do this is present in the source code:

```
pub fn ebc_cbc_detection_oracle<T: AsRef<[u8]>>(ciphertext: T, block_size: usize) -> Mode {  
    let ciphertext = ciphertext.as_ref().to_vec();  
    let similar_blocks = count_similar_blocks(ciphertext, block_size);  
    return if similar_blocks > 0 {  
        Mode::ECB  
    } else {  
        Mode::CBC  
    }  
}
```

`count_similar_blocks` is a function that simply counts similar blocks. It's actually much simpler: it's only counting repeating blocks so there's actually no concrete measure of similarity here. As we've already seen, when ECB is injected with plaintext that contains two identical blocks, it will yield the same blocks identical in ciphertext. That is a problem when it comes to cryptography because it reveals a lot about the plaintext, which is unacceptable.

There's a running saying in the cryptography community; which is that ECB lets you *see the penguin*, which is literally that, derived from a famous photo of the Linux Tux encrypted with ECB vs the same photo encrypted with another block cipher mode like CBC.

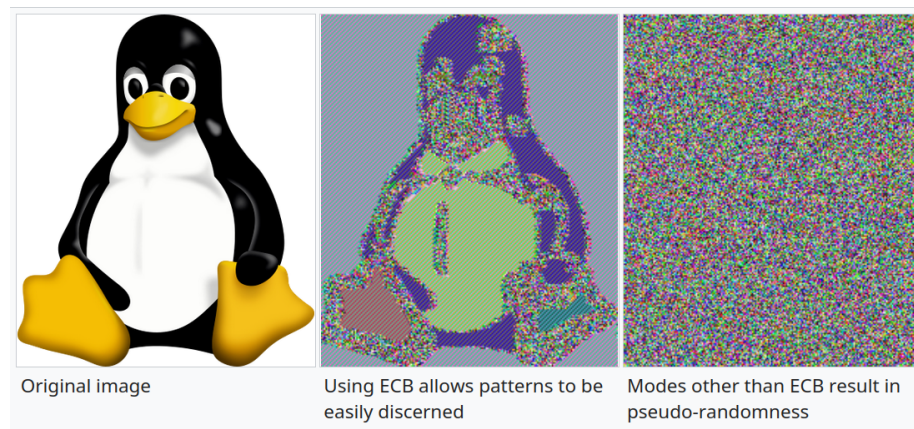


Figure 1: The famous ECB penguin

Attack 1: ECB Cut and Paste - A Replay attack To explain this attack, we will imagine a scenario in which we have a chosen plaintext attack over a client that communicates with a server securely using AES-128.

The CPA capability we have here manifests itself as the ability to inject arbitrary input into a cookie (for example), that input is then sanitized by the client before being added as a value on the cookie, for it to be encrypted with AES_ECB_128 and then sent to the server. Our only hope then of bypassing the input sanitization imposed by the client is to crack the cryptography on the wire (we suppose we have access to the channel and we can listen on the *encrypted*) communication.

When the encrypted cookie reaches the server, the server decrypts it using the symmetric key. Then, it parses the key-value pairs on the cookie to recover a specific critical field (we must note here that we don't have direct write access to this field on the client side.)

The goal of this simulation will be to demonstrate how would we go ahead about changing the critical field using only cryptography and without direct access to the field from the controlled input.

The client implementation takes in an input email, and sanitizes it from any '=' or '&' characters, it then encodes it into the form `email=<input email>&uid=<generated uid>&role=<role>` before encrypting it with AES_ECB_128 and sending it to the server. For the unprivileged user, we have `role=user` and `role=admin` for a user with high privileges. The goal of this attack is to have the client send the server the data with `role=admin` so our user can have higher privileges on the server. Again we can't just do this by injecting `email@example.com&role=admin` into the input field, because the client will clean up the '&' and '=' characters and we suppose we have no other way of bypassing this. So the only way left would be to attack the cryptography behind this: we exploit the weakness of ECB and try to craft blocks and combine and resend them to eventually have the intended payload when the server decrypts it. Evidently, we suppose we can listen in on the channel between the client and server and intercept and modify messages.

For the details, we input for example "AAAAAAAAAAadmin" padded to 26 length, it then becomes

```
email=AAAAAAAAAAadmin[ padding ]&uid=10&role=admin
      | we intercept |
      |  and save    |
      |  this block  |
```

The padding signifies padding of the block to 26 using PKCS#7 ie. padding 'admin' to 16. the block captured is a perfectly padded block that should come in the end of the intended payload. Now we have to capture the first part; for this we have x characters to play with depending on the uid field.

```
email=abc@email.com&uid=10&admin=user[ padding ]
      1           |           2           |           3           |
```

or:

```
email=abcbabc@verylongemaildomain.com&uid=10&admin=user[ padding ]
                1           |           2           |           3           |           4           |
```

etc...

We then replace the last block with our specially crafted block that we captured earlier to get the intended payload and then send it to the server, when it gets decrypted we find:

```
email=abc@email.com&uid=10&admin=admin[ padding ]
                1           |           2           |           3           |
```

Below is the part of the code that performs this attack:

```
pub fn perform_ecb_cut_and_paste() {
    // client(input_email) == encrypt(profile_for(input_email)) -> cipher
    //      server(cipher) == check_admin(parsing_routing(decrypt(cipher)))
    let cipher = Cipher::new();
    let ciphertext = client(
        &cipher,
        &std::str::from_utf8(&crate::aes::pkcs_7_padding("AAAAAAAAAadmin", 26))
            .unwrap(),
        // the 26 is put there to get the padding at exactly 16 so i can get a
        // good encrypted block of the admin padded to 16
    );
    // client sending to server, intercepting...
    let intercepted_ciphertext = ciphertext;
    // crafting input
    let admin_block = intercepted_ciphertext.chunks(16).nth(1).unwrap().to_vec();
    let ciphertext = client(&cipher, "abc@gmail.com");
    let first_block = ciphertext
        .chunks(16)
        .take(2)
        .flatten()
        .map(|x| *x)
        .collect::<Vec<u8>>();
    let payload = vec![first_block, admin_block]
        .into_iter()
        .flatten()
        .collect();
    let result = server(&cipher, payload);
    println!("admin role = {result}");
}
```

The `profile_for` is a client side function that sanitizes input from malicious characters and appends the additional info to output the formatted text under the key-value format. Whereas `parsing_routine` is server side and does quite the opposite: parses and deconstructs the decrypted cipher. `check_admin` checks if the user has admin privileges.

Attack 2: Byte at a time ECB decryption: an ECB Chosen Plaintext

Attack This attack really shows the weakness of ECB. The setup for this demonstration is like this: suppose we have a ‘partial’ chosen plaintext attack, meaning we only control part of the plaintext, which gets mixed up with another unknown secret, and gets encrypted and sent through a channel which we can sniff on. The other secret here is modeled with a secret string which gets appended to our input. Because the encryption is done using AES_ECB_128, we can use the properties of ECB to decrypt it. As we’ve already seen, the same plaintext block yields the same ciphertext block when encrypted, it’s independent of other blocks. And since our input is exactly next to the secret we want to decrypt, we can play with the input to maybe reveal characters; specifically, we can initially feed input that is exactly one byte short of the block size, and that makes the first byte of the appended text part of the ciphertext block of our input, then we can just bruteforce it by trying every possible character (256 tries). In practice, we provide a random string (that we encode in base64 for obscurity purposes) that gets appended to the input before encryption.

```
pub fn crack_a_block(oracle: &Oracle, block_size: usize) -> Vec<u8> {
    let mut cracked_bytes: Vec<u8> = Vec::new();
    // for as long as the secret string
    for byte in 0..oracle.encrypt("").len() {
        let chunk_index = byte / block_size;
        let block: Vec<u8> =
            (0..(chunk_index + 1) * block_size - 1 - cracked_bytes.len())
                .map(|_| 65_u8) // ascii('A') is 65, 65_u8, then, is the byte 'A'
                .collect::<Vec<u8>>();
        let cipherblocks_to_compare_with = oracle.encrypt(&block)
            [chunk_index * block_size..(chunk_index + 1) * block_size]
            .to_vec();
        for i in 0..=255_u8 {
            if cipherblocks_to_compare_with
                == oracle.encrypt(
                    vec![block.clone(), cracked_bytes.clone(), vec![i]]
                        .into_iter()
                        .flatten()
                        .collect::<Vec<u8>>(),
                ) [chunk_index * block_size..(chunk_index + 1) * block_size]
                .to_vec()
            {
                cracked_bytes.push(i);
                break;
            }
        }
    }
    cracked_bytes
}
```


this attack on an implementation of AES_CBC_128. This attack depends on a side-channel leak by the decryption function. The leak here is the error message that tells us whether or not the padding is valid. This is an information on the plaintext that we shouldn't have access to. As a reminder, a valid padding is N characters of the byte representation of N (with $0 < N < \text{block size} + 1$) appended to the plaintext to make the length of the block. With this knowledge, here's a simulation of this attack in Rust:

```
pub fn attack() {
    let cipher = functions::Cipher::new();
    let (ciphertext, iv) = functions::first_function(&cipher);
    let mut plaintext: Vec<char> = Vec::new();
    let mut payload: Vec<u8> = vec![iv.to_vec(), ciphertext[..16].to_vec()]
        .into_iter()
        .flatten()
        .collect::<Vec<u8>>();
    // iv block block block
    //   inter inter inter
    //   plain plain plain
    for block_index in (0..ciphertext.len()).step_by(16) {
        let mut intermediate_block = vec![0; 16];
        let mut plain_block = vec![0 as char; 16];

        if block_index != 0 {
            payload = ciphertext[block_index - 16..block_index + 16].to_vec();
        }
        for byte_index in (0..16).rev() {
            let cipher_byte = if block_index == 0 {
                // iv
                iv[byte_index]
            } else {
                // minus 16 to fall back on ciphertext..
                ciphertext[block_index + byte_index - 16]
            };
            for byte_guess in 0..=255_u8 {
                // bruteforcing to find the byte which yeilds a valid padding
                payload[byte_index] = byte_guess;
                // if padding is valid
                if functions::second_function(&cipher, &payload) {
                    // this next block here is to check for valid padding
                    // coincidences...
                    if byte_index > 0 {
                        payload[byte_index - 1] = 0xff_u8;
                        if !functions::second_function(&cipher, &payload) {
                            // in this case weve fallen on an unexpected valid
                            // padding
                        }
                    }
                }
            }
        }
    }
}
```

```

        continue;
    }
}
// intermediate_byte = 0x01 ^ byte_guess
let intermediate_byte = 16 - byte_index as u8 ^ byte_guess;
intermediate_block[byte_index] = intermediate_byte;
let plain_byte;
// plain_byte = intermediate_byte ^ prevcipherbyte[byte_index]
plain_byte = intermediate_byte ^ cipher_byte;
// push found byte
plain_block[byte_index] = plain_byte as char;
// setting the last bytes to
// 0x02 or 0x03 0x03 or 0x04 0x04 0x04 etc...
break;
}
}
for byte in byte_index..16 {
    // padding byte for next time
    let padding_byte = (16 - byte_index + 1) as u8;
    payload[byte] = padding_byte ^ intermediate_block[byte] as u8;
}
}
plain_block.iter().for_each(|&x| plaintext.push(x));
}
println!("{:?}", plaintext.iter().collect::<String>());
}

```

This attack works by bruteforce bitflipping a character in a certain block of the ciphertext; this in turn induces a bitflip in the next plaintext block when decrypted. If we corrupt, say, the last byte of the previous to last block the server will check the padding and it will be a wrong padding, we can keep flipping the bytes until we get a valid padding message, in this case, the padding is either `\x01` or `\x02\x02` or `\x03\x03\x03` etc... then we can just XOR the byte `\x01` with the byte from the ciphertext block before to get the AES decrypted ciphertext byte. We then XOR it again with the byte from the ciphertext block before to get the plaintext byte. We need at most 256 tries to find the last byte. After that we switch to finding the second to last byte and proceed with a similar approach although we have to keep in mind that we are targeting a different padding.

History of these attacks in the TLS implementations

BEAST

POODLE

References

- <https://crypto.stackexchange.com/questions/40800/is-the-padding-oracle-attack-deterministic>